



EXTENDING R WITH C++

MOTIVATION, EXAMPLES, AND CONTEXT

Dirk Eddelbuettel

20 April 2018

Debian / R Project / U of Illinois







OUTLINE

AGENDA

- (Very) Quick R Basics Reminder
- C++ in (way less than) a nutshell
- Extending R with C++ via Rpp
- A Worked Example
- A Case Study

WHY R?

PROGRAMMING WITH DATA FROM 1977 TO 2016




Thanks to John Chambers for high-resolution cover images. The publication years are, respectively, 1977, 1988, 1992, 1998, 2008 and 2016.

A SIMPLE EXAMPLE

```
xx <- faithful[, "eruptions"]
fit <- density(xx)
plot(fit)
```


A SIMPLE EXAMPLE



A SIMPLE EXAMPLE - REFINED

```
xx <- faithful[, "eruptions"]
fit1 <- density(xx)
fit2 <- replicate(10000, {
  x <- sample(xx, replace=TRUE);
  density(x, from=min(fit1$x), to=max(fit1$x))$y
})
fit3 <- apply(fit2, 1, quantile, c(0.025, 0.975))
plot(fit1, ylim=range(fit3))
polygon(c(fit1$x, rev(fit1$x)), c(fit3[1,], rev(fit3[2,])),
        col='grey', border=F)
lines(fit1)
```

A SIMPLE EXAMPLE - REFINED



So WHY R?

R enables us to

- work interactively
- explore and visualize data
- access, retrieve and/or generate data
- summarize and report into pdf, html, ...

making it the key language for statistical computing, and a preferred environment for many data analysts.

R AS CENTRAL POINT



R AS CENTRAL POINT

From any one of

- csv
- txt
- xlsx
- xml, json, ...
- web scraping, ...
- hdf5, netcdf, ...
- sas, stata, spss, ...
- various SQL + NOSQL DBs
- various binary protocols


via



into any one of

- txt
- html
- latex and pdf
- html and js
- word
- shiny
- most graphics formats
- other dashboards
- web frontends

R PER JOHN CHAMBERS (2016)



Three Principles (Section 1.1)

Object Everything that exists in R is an object.

Function Everything that happens in R is a function call.

Interface Interfaces to other software are part of R.

Three Principles (Section 1.1)


Object Everything that exists in R is an object.

Function Everything that happens in R is a function call.

Interface Interfaces to other software are part of R.

That is new. Or is it?

R PER JOHN CHAMBERS (2016)



Source: John Chamber, personal communication

INTERFACE 1976

This became the system known as “Interface”, a precursor to S and R.

C++

WHY C++?

- Asking Google leads to tens of million of hits.
- [Wikipedia](#): *C++ is a statically typed, free-form, multi-paradigm, compiled, general-purpose, powerful programming language*
- C++ is industrial-strength, vendor-independent, widely-used, and *still evolving*
- In science & research, one of the most frequently-used languages: If there is something you want to use / connect to, it probably has a C/C++ API
- As a widely used language it also has good tool support (debuggers, profilers, code analysis)

WHY C++?

Scott Meyers: *View C++ as a federation of languages*

- C provides a rich inheritance and interoperability as Unix, Windows, ... are all build on C.
- *Object-Oriented C++* (maybe just to provide endless discussions about exactly what OO is or should be)
- *Templated C++* which is mighty powerful; template meta programming unequalled in other languages.
- *The Standard Template Library (STL)* is a specific template library which is powerful but has its own conventions.
- C++11 and C++14 (and beyond) add enough to be called a fifth language.

NB: Meyers original list of four languages appeared years before C++11.

WHY C++?

- Mature yet current
- Strong performance focus:
 - *You don't pay for what you don't use*
 - *Leave no room for another language between the machine level and C++*
- Yet also powerfully abstract and high-level
- C++11, C++14, C++17, ... a big deal giving us new language features
- While there are complexities, Rcpp users are mostly shielded

C++ IN TOO LITTLE TIME

COMPILED NOT INTERPRETED

Need to compile and link

```
#include <cstdio>

int main(void) {
    printf("Hello, world!\n");
    return 0;
}
```

COMPILED NOT INTERPRETED

Or streams output rather than `printf`

```
#include <iostream>

int main(void) {
    std::cout << "Hello, world!" << std::endl;
    return 0;
}
```

COMPILED NOT INTERPRETED

`g++ -o` will compile and link

Next: an example with explicit linking of an external library.

COMPILED NOT INTERPRETED

```
#include <stdio>

#define MATHLIB_STANDALONE
#include <Rmath.h>

int main(void) {
    printf("N(0,1) 95th percentile %9.8f\n",
           qnorm(0.95, 0.0, 1.0, 1, 0));
}
```

COMPILED NOT INTERPRETED

We may need to supply:

- *header location* via `-I`,
- *library location* via `-L`,
- *library* via `-llibraryname`

```
g++ -I/usr/include -c qnorm_rmath.cpp
```

```
g++ -o qnorm_rmath qnorm_rmath.o -L/usr/lib -lRmath
```

STATICALLY TYPED

- R is dynamically typed: `x <- 3.14; x <- "foo"` is valid.
- In C++, each variable must be declared before first use.
- Common types are `int` and `long` (possibly with `unsigned`), `float` and `double`, `bool`, as well as `char`.
- No standard string type, though `std::string` is close.
- All these variables types are scalars which is fundamentally different from R where everything is a vector.
- `class` (and `struct`) allow creation of composite types; classes add behaviour to data to form `objects`.
- Variables need to be declared, cannot change

C++ IS A BETTER C

- control structures similar to what R offers: **for, while, if, switch**
- functions are similar too but note the difference in positional-only matching, also same function name but different arguments allowed in C++
- pointers and memory management: very different, but lots of issues people had with C can be avoided via STL (which is something Rcpp promotes too)
- sometimes still useful to know what a pointer is ...

OBJECT-ORIENTED

This is a second key feature of C++, and it is different from S3 and S4.

```
struct Date {  
    unsigned int year;  
    unsigned int month;  
    unsigned int day  
};  
  
struct Person {  
    char firstname[20];  
    char lastname[20];  
    struct Date birthday;  
    unsigned long id;  
};
```

OBJECT-ORIENTED

Object-orientation matches data with code operating on it:

```
class Date {  
private:  
    unsigned int year  
    unsigned int month;  
    unsigned int date;  
public:  
    void setDate(int y, int m, int d);  
    int getDay();  
    int getMonth();  
    int getYear();  
}
```

GENERIC PROGRAMMING AND THE STL

The STL promotes *generic* programming.

For example, the sequence container types **vector**, **deque**, and **list** all support

- `push_back()` to insert at the end;
- `pop_back()` to remove from the front;
- `begin()` returning an iterator to the first element;
- `end()` returning an iterator to just after the last element;
- `size()` for the number of elements;

but only **list** has `push_front()` and `pop_front()`.

Other useful containers: **set**, **multiset**, **map** and **multimap**.

GENERIC PROGRAMMING AND THE STL

Traversal of containers can be achieved via *iterators* which require suitable member functions `begin()` and `end()`:

```
std::vector<double>::const_iterator si;  
for (si=s.begin(); si != s.end(); si++)  
    std::cout << *si << std::endl;
```

GENERIC PROGRAMMING AND THE STL

Another key STL part are *algorithms*:

```
double sum = accumulate(s.begin(), s.end(), 0);
```

Some other STL algorithms are

- **find** finds the first element equal to the supplied value
- **count** counts the number of matching elements
- **transform** applies a supplied function to each element
- **for_each** sweeps over all elements, does not alter
- **inner_product** inner product of two vectors

TEMPLATE PROGRAMMING

Template programming provides a ‘language within C++’: code gets evaluated during compilation.

One of the simplest template examples is

```
template <typename T>
const T& min(const T& x, const T& y) {
    return y < x ? y : x;
}
```

This can now be used to compute the minimum between two `int` variables, or `double`, or in fact any *admissible type* providing an `operator<()` for less-than comparison.

TEMPLATE PROGRAMMING

Another template example is a class squaring its argument:

```
template <typename T>
class square : public std::unary_function<T,T> {
public:
    T operator()(T t) const {
        return t*t;
    }
};
```

which can be used along with STL algorithms:

```
transform(x.begin(), x.end(), square);
```

FURTHER READING

Books by Meyers are excellent

I also like the (free) [C++ Annotations](#)

C++ FAQ

Resources on StackOverflow such as

- [general info](#) and its links, eg
- [booklist](#)

DEBUGGING

Some tips:

- Generally painful, old-school `printf()` still pervasive
- Debuggers go along with compilers: `gdb` for `gcc` and `g++`; `lldb` for the `clang` / `llvm` family
- Extra tools such as `valgrind` helpful for memory debugging
- “Sanitizer” (ASAN/UBSAN) in newer versions of `g++` and `clang++`

EXTENDING R WITH C++

Three key functions

- `evalCpp()`
- `sourceCpp()`
- `cppFunction()`

BASIC USAGE: EVALCPP()

`evalCpp()` evaluates a single C++ expression. Includes and dependencies can be declared.

This allows us to quickly check C++ constructs.

```
library(Rcpp)
evalCpp("2 + 2")      # simple test

## [1] 4

evalCpp("std::numeric_limits<double>::max()")

## [1] 1.797693e+308
```

BASIC USAGE: CPPFUNCTION()

cppFunction() creates, compiles and links a C++ file, and creates an R function to access it.

```
cppFunction("  
  int simpleExample() {  
    int x = 10;  
    return x;  
}  
simpleExample() # same identifier as C++ function
```

BASIC USAGE: CPPFUNCTION()

cppFunction() creates, compiles and links a C++ file, and creates an R function to access it.

```
cppFunction("  
  int exampleCpp11() {  
    auto x = 10;  
    return x;  
}", plugins=c("cpp11"))  
exampleCpp11() # same identifier as C++ function
```


BASIC USAGE: SOURCECPP()

`sourceCpp()` is the actual workhorse behind `evalCpp()` and `cppFunction()`. It is described in more detail in the package vignette [Rcpp-attributes](#).

`sourceCpp()` builds on and extends `cxxfunction()` from package `inline`, but provides even more ease-of-use, control and helpers – freeing us from boilerplate scaffolding.

A key feature are the plugins and dependency options: other packages can provide a plugin to supply require compile-time parameters (cf `RcppArmadillo`, `RcppEigen`, `RcppGSL`).

BASIC UAGE: RSTUDIO



BASIC UAGE: RSTUDIO (CONT'ED)

The following file gets created:

```
#include <Rcpp.h>
using namespace Rcpp;

// This is a simple example of exporting a C++ function to R. You can
// source this function into an R session using the Rcpp::sourceCpp
// function (or via the Source button on the editor toolbar). ...

// [[Rcpp::export]]
NumericVector timesTwo(NumericVector x) { return x * 2; }

// You can include R code blocks in C++ files processed with sourceCpp
// (useful for testing and development). The R code will be automatically
// run after the compilation.

/** R
timesTwo(42)
*/
```

So what just happened?

- We defined a simple C++ function
- It operates on a numeric vector argument
- We asked Rcpp to ‘source it’ for us
- Behind the scenes Rcpp creates a wrapper
- Rcpp then compiles, links, and loads the wrapper
- The function is available in R under its C++ name

BASIC USAGE: PACKAGES


Package are *the* standard unit of R code organization.

Creating packages with Rcpp is easy; an empty one to work from can be created by `Rcpp.package.skeleton()`

The vignette [Rcpp-packages](#) has fuller details.

As of April 2018, there are over 1300 packages on CRAN which use Rcpp, and almost 100 more on BioConductor — with working, tested, and reviewed examples.

PACKAGES AND RCPP



PACKAGES AND RCPP

`Rcpp.package.skeleton()` and its derivatives. e.g.

`RcppArmadillo.package.skeleton()` create working packages.

```
// another simple example: outer product of a vector,
// returning a matrix
//
// [[Rcpp::export]]
arma::mat rcpparma_outerproduct(const arma::colvec & x) {
    arma::mat m = x * x.t();
    return m;
}

// and the inner product returns a scalar
//
// [[Rcpp::export]]
double rcpparma_innerproduct(const arma::colvec & x) {
    double v = arma::as_scalar(x.t() * x);
    return v;
}
```

NICE, BUT DOES IT *REALLY WORK?*

SIMPLE EXAMPLE

Something self-contained

- Let's talk random numbers!
- We'll look at a quick generator
- And wrap it in plain C / C++

RANDOM NUMBER

<|

< PREV

RANDOM

NEXT >

|>

```
int getRandomNumber()
{
    return 4; // chosen by fair dice roll.
               // guaranteed to be random.
}
```

<|

< PREV

RANDOM

NEXT >

|>

PERMANENT LINK TO THIS COMIC: [HTTPS://XKCD.COM/221/](https://xkcd.com/221/)

XKCDRNG.H

```
// cf https://xkcd.com/221/
//
//      "RFC 1149.5 specifies 4 as the "
//      "standard IEEE-vetted random number."
//
int getRandomNumber()
{
    return 4; // chosen by fair dice roll
              // guaranteed to be random
}
```

GETXKCDRNGDRAW()

```
#include <Rcpp.h>
#include <xkcdRng.h>

// [[Rcpp::export]]
int getXkcdRngDraw() {
    return getRandomNumber();
}
```

PACKAGE

The screenshot shows the RStudio interface with the following details:

- File Menu:** File, Edit, Code, View, Plots, Session, Build, Debug, Profile, Tools, Help.
- Code Editor:** Shows code in `zdRng.h` and `getXkcdRngDraw.cpp`. The `getXkcdRngDraw()` function is annotated with `[[Rcpp::export]]`.
- Environment Tab:** Shows the command line output of the `roxygenize` and `R CMD INSTALL` commands.
- Console Tab:** Shows the directory structure of the package (`DESCRIPTION`, `man`, `src` containing `RcppExports.cpp` and `xkcdRng.h`).
- Terminal Tab:** Shows the command `tree` output.
- Build Tab:** Shows the build log for the package.
- Files Tab:** Shows the contents of the `git` folder, including `gitignore`, `history`, `DESCRIPTION`, and `man` directories.

```
#include <Rcpp.h>
#include <xkcdRng.h>
// [[Rcpp::export]]
int getXkcdRngDraw() {
  return getRandomNumber();
}

** roxygen2::roxygenize("./", roclets=c("rd"))

First time using roxygen2. Upgrading automatically...
Documentation completed

==> R CMD INSTALL --no-multiarch --with-keep.source samplexkcdrng

ccache g++ -I/usr/share/R/include -DNDEBUG -I"/usr/local/lib/R/site-library/Rcpp/include" -I. -fPIC -g -O3 -Wall -pipe -Wno-unused -march=native -c RcppExports.cpp -o RcppExports.o
* installing to library '/usr/local/lib/R/site-library'
* installing "source" package 'samplexkcdrng' ...
** libs
ccache g++ -I/usr/share/R/include -DNDEBUG -I"/usr/local/lib/R/site-library/Rcpp/include" -I. -fPIC -g -O3 -Wall -pipe -Wno-unused -march=native -c getXkcdRngDraw.cpp -o getXkcdRngDraw.o
ccache g++ -Wl,-S -shared -L/usr/lib/R/lib -Wl,-Bsymbolic-functions -Wl,-z,relro -o samplexkcdrng.so RcppExports.o getXkcdRngDraw.o -L/usr/lib/R/lib -L/usr/lib
installing to /usr/local/lib/R/site-library/samplexkcdrng/libs
** R
** preparing package for lazy loading
** help
*** installing help indices
** building package indices
** testing if installed package can be loaded
* DONE (samplexkcdrng)
```

WHAT DID WE DO?

- An unmodified piece of C / C++ code
- A simple interface function
- Rcpp does the rest

A CASE STUDY

A BENCHMARK COMPARISON

A recent blogpost on “finding a needle in a haystack” has a nice story:

```
options(width=50)
set.seed(1)
haystack <- sample(0:12, size = 2000, replace = TRUE)
needle    <- c(2L, 10L, 8L)
haystack[1:60]
```

```
##  [1]  3   4   7  11   2  11  12   8   8   0   2   2   8   4  10
## [16]  6   9  12   4  10  12   2   8   1   3   5   0   4  11   4
## [31]  6   7   6   2  10   8  10   1   9   5  10   8  10   7   6
## [46] 10   0   6   9   9   6  11   5   3   0   1   4   6   8   5
```

A BENCHMARK COMPARISON

A recent blogpost on “finding a needle in a haystack” has a nice story:

```
options(width=50)
set.seed(1)
haystack <- sample(0:12, size = 2000, replace = TRUE)
needle    <- c(2L, 10L, 8L)
haystack[1:60]

## [1]  3  4  7 11  2 11 12  8  8  0  2  2  8  4 10
## [16]  6  9 12  4 10 12  2  8  1  3  5  0  4 11  4
## [31]  6  7  6  2 10  8 10  1  9  5 10  8 10  7  6
## [46] 10  0  6  9  9  6 11  5  3  0  1  4  6  8  5
```

FIRST CANDIDATE

```
forloop_find <- function(needle, haystack) {  
  n <- length(needle) - 1L  
  for (i in seq(haystack)) {  
    if (identical(haystack[i:(i+n)], needle)) {  
      return(i)  
    }  
  }  
}  
  
forloop_find(needle, haystack)
```

```
## [1] 34
```

SECOND CANDIDATE

```
lead_find <- function(needle, haystack) {  
  v <- haystack == needle[1]  
  for (i in seq(2, length(needle))) {  
    v <- v +  
      (dplyr::lead(haystack, i-1L) == needle[i])  
  }  
  which(v == length(needle))[1L]  
}  
lead_find(needle, haystack)  
  
## [1] 34
```

THIRD CANDIDATE

```
shift_find <- function(needle, haystack) {  
    shifted_haystack <-  
        data.table::shift(haystack, type='lead',  
                           0:(length(needle)-1))  
  
    v <- Map('==' , shifted_haystack, needle)  
    v <- Reduce('+', v)  
    which(v == length(needle))[1]  
}  
shift_find(needle, haystack)  
  
## [1] 34
```

FOURTH CANDIDATE

```
Rcpp::cppFunction('int rcpp_find(NumericVector needle,  
                           NumericVector haystack) {  
    int nlen = needle.size(), hlen = haystack.size(), j;  
    for (int i = 0; i < (hlen - nlen); i++) {  
        for (j = 0; j < nlen; j++) {  
            if (needle[j] != haystack[i + j]) break;  
        }  
        if (j == nlen) return(i+1);  
    }  
    return(0);  
}' )  
rcpp_find(needle, haystack)
```

[1] 34

U Illinois Stat 385 Guest Lecture

65/71

FIFTH CANDIDATE

```
Rcpp::cppFunction('
int idiomaticrcpp_find(NumericVector needle,
                           NumericVector haystack) {
    NumericVector::iterator it;
    it = std::search(haystack.begin(), haystack.end(),
                     needle.begin(), needle.end());
    int pos = it - haystack.begin() + 1;
    if (pos > haystack.size()) pos = -1;
    return(pos);
}')
idiomaticrcpp_find(needle, haystack)
```

```
## [1] 34
```


SHOOTOUT

```
R> res <- microbenchmark::microbenchmark(...) # not shown
R> res
Unit: microseconds

      expr      min       lq      mean     median       uq      max neval
rollapply_find(needle, haystack) 5829.484 6355.7290 6918.75051 6719.5825 7114.4770 37348.915 1000
  vapply_find(needle, haystack) 1230.373 1338.3275 1519.16471 1419.7170 1491.9485 31687.188 1000
    grep_find(needle, haystack) 535.059 556.3545 592.86697 572.5000 597.6460 1396.680 1000
forloop_find(needle, haystack) 169.751 189.5370 213.39386 200.6070 210.2785 1151.483 1000
    lead_find(needle, haystack) 47.571 55.4575 61.54025 59.2370 62.3445 331.499 1000
    shift_find(needle, haystack) 37.939 47.5780 55.14043 52.3475 58.5400 268.533 1000
      pile_find(needle, haystack) 13.883 15.7375 17.10174 16.5590 17.4175 45.757 1000
    sieved_find(needle, haystack) 7.587 9.4575 10.82973 10.4355 11.3620 29.978 1000
  boyer_moore(needle, haystack) 3.197 4.7035 6.06770 5.6540 6.5950 89.414 1000
      rcpp_find(needle, haystack) 2.579 3.6805 4.86756 4.5465 5.3570 27.765 1000
idiomaticrcpp_find(needle, haystack) 2.183 3.5230 4.62004 4.3555 5.1945 24.235 1000
R>
```

SHOOTOUT

```
ggplot2:::autoplot(res)
```



CONCLUSION

Takeaways on *Extending R with C++*

- clearly possible as the tooling helps greatly
- natural as interfaces are a normal part of R
- not too hard, though balancing two languages
- rewarding in terms of performance
- always measure and profile

THANK YOU!

slides <http://dirk.eddelbuettel.com/presentations/>

web <http://dirk.eddelbuettel.com/>

mail dirk@eddelbuettel.com

github [@eddelbuettel](https://github.com/eddelbuettel)

twitter [@eddelbuettel](https://twitter.com/eddelbuettel)